Recovering Sparse Neural Connectivity from Partial Measurements: A Covariance-Based Approach with Granger-Causality Refinement

Quilee Simeon
Massachusetts Institute of Technology
qsimeon@mit.edu

Abstract

Inferring the connectivity of neural circuits from incomplete observations is a fundamental challenge in neuroscience. We present a covariance-based method for recovering the weight matrix of a recurrent neural network from partial measurements across multiple recording sessions, accumulating pairwise covariance estimates from sessions with overlapping neuron subsets. A Granger-causality refinement step enforces biological constraints via projected gradient descent, achieving $r = 0.91$ correlation with true connectivity and 82% improvement over chance using only 66% neuron coverage per session. Through systematic experiments on synthetic networks modeling small brain circuits, we characterize a fundamental *control-estimation tradeoff*: stimulation aids identifiability but disrupts intrinsic dynamics, with the optimal level depending on measurement density. We discover that the "incorrect" linear approximation acts as implicit regularization—matching or outperforming the oracle estimator with known nonlinearity at all tested regimes—a concrete instance of James–Stein shrinkage. Error decomposition reveals that correlation from central pattern generators, not model mismatch, is the dominant error bottleneck ($3.4\times$ larger), identifying intrinsic dynamics modeling as the key target for further improvement.

1 Introduction

A central goal in systems neuroscience is to map the connectivity of neural circuits from functional measurements [Power et al., 2011]. In small brain circuits—such as *C. elegans* [White et al., 1986, Cook et al., 2019], larval zebrafish [Ahrens et al., 2013], and cortical organoids [Quadrato et al., 2017]—the number of neurons is tractable (10^2–10^4) but simultaneous recording of the entire circuit remains difficult. This creates a fundamental inverse problem: given partial, noisy measurements of neural dynamics across multiple sessions, can we recover the full connectivity matrix?

This problem embodies a *control-estimation tradeoff* related to classical notions of persistent excitation in system identification [Ljung, 1998]: stimulating neurons with random perturbations aids connectivity estimation by exciting network modes, but disrupts the intrinsic dynamics that make the circuit biologically relevant. An ideal experimental protocol must balance identifiability against preservation of natural dynamics. Furthermore, real circuits possess intrinsic oscillatory patterns—central pattern generators (CPGs) [Marder and Bucher, 2001]—whose detailed structure is unknown to the experimenter. These represent *unmodeled autonomous dynamics* that influence the observed activity.

Preprint. Under review.

We formalize this as a matrix recovery problem from partial observations. Consider a network of N neurons with connectivity matrix $W \in \mathbb{R}^{N\times N}$, governed by discrete-time recurrent dynamics [Sussillo and Abbott, 2009], where W_{ij} represents the synaptic weight from neuron j to neuron i. In each recording session, we observe a subset of neurons and stimulate a (possibly different) subset. Our key insight is that pairwise covariance statistics can be accumulated across sessions: whenever neurons i and j are simultaneously observed, their covariance contributes to estimating the (i, j) block of W.

Contributions. Our method is general, applying to any network where the same connectivity is measured across multiple partial-observation sessions.

- We derive a covariance-based estimator $\hat{W} = \Sigma_{x_{t+1},x_t}\Sigma^{-1}_{x_t,x_t}$ for connectivity recovery from partial measurements and provide identifiability conditions.
- We introduce a Granger-causality-inspired [Granger, 1969] refinement via projected gradient descent that enforces biological constraints (sparsity, non-negativity, no autapses).
- We discover an *implicit regularization* effect: the "incorrect" linear approximation yields a better-conditioned estimator than the oracle with known nonlinearity, analogous to James–Stein shrinkage [James and Stein, 1961].
- We systematically characterize the tradeoff between stimulation strength and measurement density, revealing that optimal stimulation depends on the fraction of observed neurons.

2 Problem Formulation

2.1 Dynamical System Model

We model the neural circuit as a discrete-time recurrent dynamical system:

$$x_{t+1} = W\phi(x_t) + b_t \tag{1}$$

where $x_t \in \mathbb{R}^N$ is the neural state, $W \in \mathbb{R}^{N\times N}$ is the connectivity matrix, $\phi : \mathbb{R} \to \mathbb{R}$ is an element-wise nonlinearity, and $b_t \in \mathbb{R}^N$ is the total input comprising both extrinsic stimulation and intrinsic central pattern generator (CPG) drive. This arises from discretizing a continuous-time neural circuit $\tau\, d\mathbf{x}/dt = -\tilde{M}\phi(\mathbf{x}) + \tilde{\mathbf{b}}$ with time constant τ and coupling matrix $\tilde{M}$, yielding $W = I - (\Delta t/\tau)\tilde{M}$.

Network topology and weights. We impose $W_{ii} = 0$ for all i (no autapses/self-connections), a known structural constraint in *C. elegans* and many neural circuits [White et al., 1986, Cook et al., 2019]. The off-diagonal entries of W are sampled as a random sparse directed graph: we find the minimal density p such that the digraph is strongly connected (matching the sparse, connected topology observed in real connectomes), then draw non-zero entries from $\mathrm{Unif}(0, 1)$ and scale so that the spectral radius $\rho(W) \le 1$. We focus on *non-negative* weights ($W_{ij} \ge 0$). This choice is motivated by the measurement modality: electron microscopy (EM) connectomics—the primary source of ground-truth wiring data—measures synapse *size* (contact area, vesicle count) but not synapse *sign* (excitatory vs. inhibitory), yielding an inherently unsigned weight matrix. In our formulation, inhibitory effects arise from the *sign of the activity*: a positive weight applied to negative (hyperpolarized) presynaptic activity drives the postsynaptic neuron downward. This is a dual of the standard neuroscience convention (non-negative firing rates with signed weights), and is functionally equivalent for the covariance estimator, which assumes no sign structure on W. By the Perron–Frobenius theorem, non-negative irreducible matrices have a unique maximal real eigenvalue with a positive eigenvector, which simplifies spectral radius control and ensures stable dynamics. The method generalizes to mixed-sign weights (see Section 6).

Nonlinearity. We require the element-wise nonlinearity ϕ to be *1-Lipschitz*: $|\phi'(x)| \le 1$ for all x. Combined with $\rho(W) \le 1$, this guarantees contractivity of the autonomous dynamics. Our default is $\phi = \tanh$, which is 1-Lipschitz, bounded, and odd-symmetric ($\tanh(-x) = -\tanh(x)$, so $\tanh(0) = 0$); these additional properties enable exact error analysis via the Stein–Price identity (Section 3). To assess robustness, we test three additional nonlinearities in E5 (Section 4): *identity* ($\phi(x) = x$, 1-Lipschitz but unbounded), *ReLU* ($\phi(x) = \max(0, x)$, 1-Lipschitz but neither odd nor bounded), and *sigmoid* ($\phi(x) = 1/(1 + e^{-x})$, 1-Lipschitz and bounded but not odd, with

$\phi(0) = 0.5 \neq 0$). This deliberately includes nonlinearities that violate our default's symmetry and boundedness, testing how gracefully the estimator degrades.

Stability. For stability, we require $\rho(W) \leq 1$, where $\rho(\cdot)$ denotes the spectral radius. Combined with the 1-Lipschitz property of ϕ, this ensures bounded autonomous dynamics ($b_t = 0$): since $\|\phi(x)\| \leq \|x\|$ and $\rho(W) \leq 1$, the map $x \mapsto W\phi(x)$ is non-expansive. The spectral radius is scaled to $\rho(W) \approx 1.0$ (edge of stability), producing marginally stable dynamics that are neither trivially decaying nor divergent. Crucially, with $\rho(W) \leq 1$ and 1-Lipschitz ϕ satisfying $|\phi(x)| < |x|$ for $x \neq 0$ (as with $\tanh$), the autonomous system ($b_t = 0$) *contracts to the origin*: the network falls silent without input. This motivates the central pattern generator (CPG) component of the input b_t: a subset of neurons possess autonomous oscillatory drive—analogous to cardiac pacemaker cells—that sustains network activity in the absence of external stimulation [Marder and Bucher, 2001]. The input thus decomposes as $b_t = b_t^{\text{stim}} + b_t^{\text{CPG}}$, where b_t^{stim} is extrinsic i.i.d. Gaussian noise (independent of x_t) and $b_t^{\text{CPG}} = \text{cpg}(x_t)$ is state-dependent intrinsic drive with $\text{Cov}(b_t^{\text{CPG}}, x_t) \neq 0$. This decomposition is critical: the extrinsic component satisfies our independence assumption, while the CPG component violates it and dominates the estimation error (Section 6).

2.2 Observation Model

In session k, we observe a subset $\mathcal{M}_k \subseteq \{1, \ldots, N\}$ of neurons and stimulate a subset $\mathcal{S}_k$. The observed activity is:

$$y_t^{(k)} = P_{\mathcal{M}_k} x_t + \epsilon_t, \qquad \epsilon_t \sim \mathcal{N}(0, \sigma_\epsilon^2 I) \tag{2}$$

where $P_{\mathcal{M}_k}$ is a projection onto the observed neurons and ϵ_t is i.i.d. measurement noise. Our experiments test both noiseless ($\sigma_\epsilon = 0$) and noisy ($\sigma_\epsilon > 0$) regimes; the method is robust to moderate noise (Section 6).

2.3 The Estimation Problem

Given K sessions with observations $\{y_t^{(k)}\}_{t=1}^{T_k}$ for $k = 1, \ldots, K$, recover W.

3 Methods

3.1 Covariance-Based Estimator

Under the linear approximation $\phi(x) \approx x$ (valid when $\|x\|$ is small), Eq. (1) simplifies to $x_{t+1} \approx Wx_t + b_t$. Assuming statistical independence between inputs b_t and states x_t (an approximation we revisit critically in Section 6), the cross-covariance of the *true* states satisfies:

$$\Sigma_{x_{t+1},x_t} = W\Sigma_{x_t,x_t} + \underbrace{\Sigma_{b_t,x_t}}_{\approx 0} \tag{3}$$

In the noiseless case ($\sigma_\epsilon = 0$, so $y_t = x_t$), this directly yields the estimator $\hat{W} = \Sigma_{y_{t+1},y_t}\Sigma_{y_t,y_t}^{-1}$. When observation noise is present ($\sigma_\epsilon > 0$), the observed covariances differ from the true state covariances. Since ϵ_t is i.i.d. and independent of x_t:

$$\Sigma_{y_t,y_t} = \Sigma_{x_t,x_t} + \sigma_\epsilon^2 I \tag{4}$$

$$\Sigma_{y_{t+1},y_t} = \Sigma_{x_{t+1},x_t} \quad \text{(noise at } t{+}1 \text{ independent of } y_t) \tag{5}$$

The estimator computed from observed data is therefore:

$$\hat{W} = \Sigma_{y_{t+1},y_t}\big(\Sigma_{y_t,y_t}\big)^{-1} = W\Sigma_{x,x}\big(\Sigma_{x,x} + \sigma_\epsilon^2 I\big)^{-1} \tag{6}$$

This is a *ridge-regularized* version of the true W: observation noise shrinks the estimate toward zero, with σ_ϵ^2 playing the role of the ridge parameter. For small noise ($\sigma_\epsilon \ll \sigma_{\min}(\Sigma_{x,x})$), the effect is negligible; for large noise, the estimate is over-regularized.

No-autapse prior. Before any further refinement, we enforce the structural constraint $\hat{W}_{ii} = 0$ for all i. The diagonal of the raw estimate $\hat{\Sigma}_{y_{t+1},y_t}\hat{\Sigma}_{y_t,y_t}^{-1}$ is dominated by single-neuron autocorrelation, which does not reflect genuine self-connectivity (since we imposed $W_{ii} = 0$ by construction). Zeroing the diagonal removes this artifact and is the single most impactful step in the pipeline: in ablation experiments (E4, Section 4), it accounts for the majority of the improvement over the chance baseline.

Accumulation across sessions. For each pair (i, j), we compute empirical covariances from the *observed* data y_t only in sessions where both neurons i and j are simultaneously measured. Let $S_{ij}^{(k)} = \mathbf{1}[i \in \mathcal{M}_k \wedge j \in \mathcal{M}_k]$ indicate co-measurement. The accumulated covariance is:

$$\hat{\Sigma}_{ij} = \frac{\sum_{k=1}^{K} S_{ij}^{(k)} \hat{\Sigma}_{ij}^{(k)}}{\sum_{k=1}^{K} S_{ij}^{(k)}} \tag{7}$$

where $\hat{\Sigma}_{ij}^{(k)}$ is the empirical covariance of the observed y_t from session k.

Error analysis. The estimation error decomposes into two terms:

$$\hat{W} - W = \underbrace{W\left[\Sigma_{\phi(x),x} - \Sigma_{x,x}\right]\Sigma_{x,x}^{-1}}_{E_1\text{: model mismatch}} + \underbrace{\Sigma_{b,x}\Sigma_{x,x}^{-1}}_{E_2\text{: input correlation}} \tag{8}$$

For $\phi = \tanh$ and Gaussian states $x \sim \mathcal{N}(0, \Sigma)$, the Stein–Price identity [Price, 1958, Stein, 1981] gives an *exact* characterization. Since $\mathbb{E}[\tanh(x_i)x_j] = \Sigma_{ij}\,\mathbb{E}[\mathrm{sech}^2(x_i)]$, we have:

$$\Sigma_{\phi(x),x} = D\Sigma_{x,x}, \qquad D = \mathrm{diag}\big(\mathbb{E}[\mathrm{sech}^2(x_i)]\big) \tag{9}$$

where $0 < D_{ii} < 1$ for $\sigma_i > 0$. Since $E_1 = W(D - I)\Sigma_{x,x}\Sigma_{x,x}^{-1} = W(D - I)$, the mismatch error simplifies to $\|E_1\|_F = \|W(D - I)\|_F \leq \|W\|_F \max_i(1 - \mathbb{E}[\mathrm{sech}^2(x_i)])$. For small state variance, $1 - \mathbb{E}[\mathrm{sech}^2(x_i)] \approx \sigma_i^2$, so the error is $O(\max_i \sigma_i^2)$—substantially tighter than generic Lipschitz bounds. Crucially, $\|\Sigma_{x,x}^{-1}\|_2$ amplifies both error sources: ill-conditioned covariance (from sparse measurement) degrades recovery even with perfect model knowledge.

3.2 Granger-Causality Refinement

The raw estimator $\hat{W}$ may violate biological constraints. We refine it using projected gradient descent that enforces:

1. **No self-connections**: $W_{ii} = 0$ for all i (no autapses).
2. **Granger non-causality**: $W_{ij} = 0$ where $\Sigma_{x_t,x_t}(i, j) > \Sigma_{x_{t+1},x_t}(i, j)$, indicating that the lagged covariance does not exceed the contemporaneous covariance.
3. **Non-negativity**: $W_{ij} \geq 0$ (when applicable, e.g., for excitatory-only networks).

The optimization solves:

$$\min_A \|A - \hat{\Sigma}_{x_{t+1},x_t}\|_F^2 \quad \text{s.t.} \quad A\hat{\Sigma}_{x_t,x_t}^{-1} \in \mathcal{C} \tag{10}$$

where $\mathcal{C}$ is the constraint set. We alternate between gradient steps on A and projection of $W = A\hat{\Sigma}_{x_t,x_t}^{-1}$ onto $\mathcal{C}$.

3.3 Full Pipeline

The complete method is summarized in Algorithm 1. Figure 1 illustrates the partial observation model, and Figure 2 shows the full estimation pipeline.

4 Experiments

4.1 Experimental Setup

We generate synthetic networks modeling small brain circuits:

- Random directed graphs with non-negative weights and spectral radius $\rho(W) \leq 1$
- Intrinsic dynamics driven by central pattern generators (CPGs) modeled as chaotic reservoir networks

Algorithm 1 Covariance Accumulation + Granger Refinement

Require: Sessions $\{(y_t^{(k)}, \mathcal{M}_k)\}_{k=1}^K$
Ensure: Estimated connectivity $\hat{W}$

```
Initialize Σ̂_{y_{t+1},y_t} ← 0, Σ̂_{y_t,y_t} ← 0, C ← 0 {accumulators}
for k = 1, ..., K do
    Compute session covariances from observed data {y_t^(k)}_{t=1}^{T_k}
    S ← 1_{M_k} 1_{M_k}^T {co-measurement mask}
    Σ̂_{y_{t+1},y_t} += S ⊙ Σ̂^(k)_{y_{t+1},y_t};   Σ̂_{y_t,y_t} += S ⊙ Σ̂^(k)_{y_t,y_t}
    C += S
end for
Σ̂_{y_{t+1},y_t} ← Σ̂_{y_{t+1},y_t} ⊘ C;   Σ̂_{y_t,y_t} ← Σ̂_{y_t,y_t} ⊘ C {average}
Ŵ_approx ← Σ̂_{y_{t+1},y_t} · Σ̂†_{y_t,y_t} {ridge-regularized estimator, Eq. (6)}
diag(Ŵ_approx) ← 0 {no-autapse structural prior}
Ŵ ← GrangerRefine(Ŵ_approx, Σ̂_{y_t,y_t}, Σ̂_{y_{t+1},y_t}, C) {Eq. (10)}
return Ŵ
```

- Extrinsic stimulation as *independent* i.i.d. Gaussian noise applied to each sensor node at each timestep (the most favorable case for identifiability; correlated or shared stimulation would reduce the rank of the input covariance)

Each node in the network is characterized by three (non-exclusive) Boolean roles: whether it is a *CPG node* (receives autonomous intrinsic drive), a *sensor node* (receives extrinsic stimulation), and a *measured node* (its activity is recorded). These roles are assigned randomly per session, modeling the experimental reality that different recording sessions observe and stimulate different neuron subsets.

Our experiments systematically vary the parameters in Table 1 while holding others at baseline values ($N = 12$, $T = 900$, 8 measured neurons, 4 sensors, 4 CPG nodes, $\sigma = 1.0$, $\phi = \tanh$). Since $N = 12$, converting fractions to integer neuron counts introduces discretization: the baseline uses $8/12 \approx 66\%$ measured, $4/12 \approx 33\%$ sensors and CPGs. In sweep experiments (E2, E3), we report the target fraction (e.g., "33%") while noting that the actual count $\lfloor p \cdot N \rfloor$ may differ slightly (e.g., $\lfloor 0.33 \times 12 \rfloor = 3$ neurons $= 25\%$). Each experiment uses 17–30 repetitions (distinct random network topologies) with $K = 50$ recording sessions per topology, each sharing the same W but drawing independent random measurement masks $\mathcal{M}_k$ and sensor subsets $\mathcal{S}_k$ (parallelized via joblib). We report median Frobenius distance $\|W - \hat{W}\|_F / N$ with 95% bootstrap confidence intervals ($n_{\text{resamples}} = 1000$).

Table 1: Experimental design: each experiment varies one or two parameters while holding others at baseline.

Exp	Parameter swept	Values	Question
E1	Network size N, duration T	$N \in \{8, 12, 30\}$, $T \in \{100, 500, 1000\}$	How does recovery scale?
E2	Measurement fraction	33–100%	How much observation is needed?
E3	Stim gain $\times$ measurement	$\sigma \in [0, 2]$, meas $\in \{33, 66, 100\%\}$	Is there a stim-measurement interaction?
E4	Method ablation	Chance $\rightarrow$ Estimate $\rightarrow$ Granger	What does each step add?
E5	Nonlinearity ϕ	tanh, ReLU, identity, sigmoid	How robust to model mismatch?
E6	Oracle vs. approximation	$\sigma \in [0, 5]$	Does the oracle ever win?
E7	Sensor fraction	8–100%	How many neurons to stimulate?

The subsections below present experiments in narrative order rather than numerical order: we first establish baseline scaling (E1) and ablate method components (E4), then examine the key experimental design tradeoffs—stimulation intensity (E3) and measurement density (E2)—and finally test robustness to nonlinearity (E5), the oracle comparison (E6), and sensor coverage (E7).

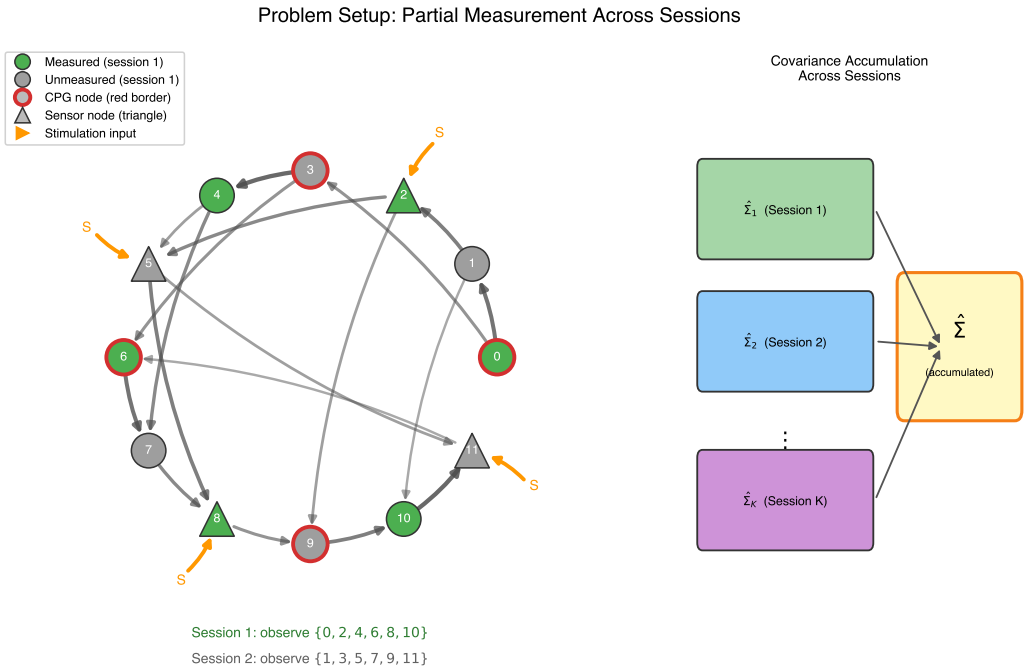

Figure 1: **Problem setup: partial measurement across recording sessions. (Left)** A network of $N = 12$ neurons with directed weighted connectivity W. In Session 1, a subset of neurons (green circles) is measured while the rest (gray) are unobserved. Nodes with red borders are central pattern generator (CPG) drivers; triangle-shaped nodes are sensor neurons receiving extrinsic stimulation (orange arrows labeled "S"). **(Right)** Covariance accumulation across K sessions: each session yields a partial covariance matrix $\hat{\Sigma}_k$ covering only co-observed neuron pairs. These are averaged element-wise to reconstruct the full accumulated covariance $\hat{\Sigma}$, from which W is estimated via Eq. (6).

4.2 Baseline Recovery (E1)

Table 2 shows recovery error across network sizes $N \in \{8, 12, 30\}$ and recording durations $T \in \{100, 500, 1000\}$ with 66% measurement. The covariance estimator (with diagonal zeroed) consistently outperforms the chance baseline, achieving median errors of 0.06–0.26 vs. ~ 0.54 for chance. Note that the estimate already enforces the no-autapse prior ($W_{ii} = 0$), which removes the autocorrelation-dominated diagonal. For $N = 30$ with $T = 1000$, the method achieves excellent recovery: 0.053 (Granger-refined), a 91% improvement over chance (Figure 3).

Table 2: Median recovery error (Frobenius / N) across network sizes and durations. Bold indicates best method per row. Gr. Improv. is the relative reduction from Estimate to Granger: $(E_{\text{est}} - E_{\text{gr}})/E_{\text{est}}$.

N	T	Chance	Estimate	Granger	Gr. Improv.
8	100	0.540	0.172	**0.156**	9%
8	500	0.540	0.143	**0.134**	6%
8	1000	0.540	0.131	**0.127**	3%
12	100	0.542	0.263	**0.163**	38%
12	500	0.542	0.109	**0.100**	8%
12	1000	0.542	0.115	**0.102**	11%
30	100	0.561	0.169	**0.098**	42%
30	500	0.561	0.063	**0.054**	13%
30	1000	0.561	0.059	**0.053**	10%

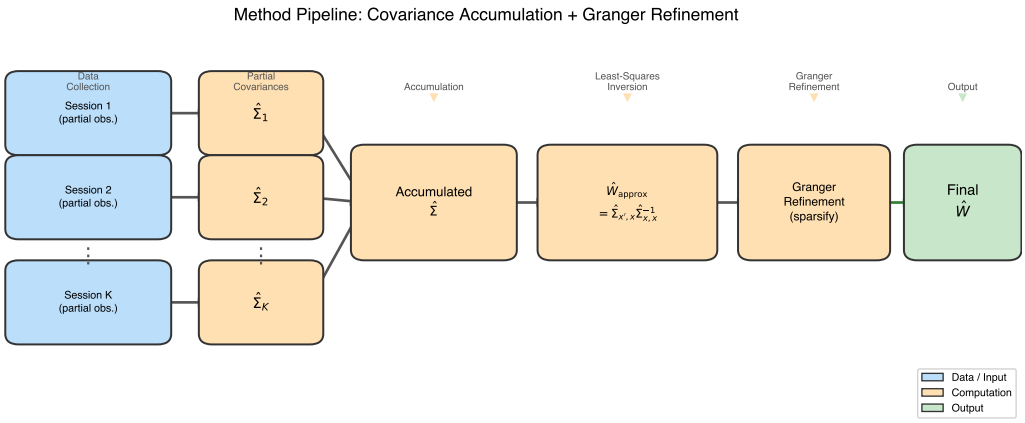

Figure 2: **Method pipeline.** Data from K recording sessions yield partial covariance matrices, which are accumulated element-wise into a full covariance estimate $\hat{\Sigma}$. The connectivity matrix is estimated via least-squares inversion, then refined by projected gradient descent enforcing biological constraints (Algorithm 1).

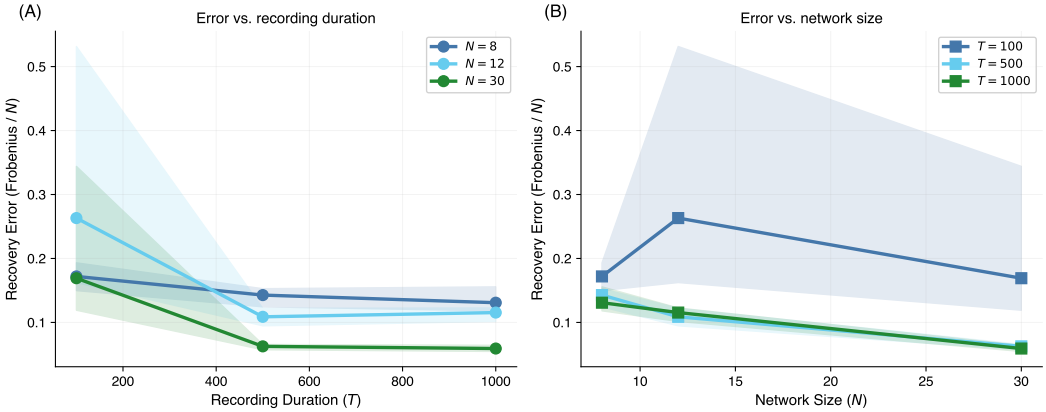

Figure 3: **Scaling properties of the covariance estimator. (Left)** Median recovery error (Frobenius distance / N) vs. recording duration T for three network sizes ($N \in \{8, 12, 30\}$) with 66% measurement. Shaded regions show 95% bootstrap confidence intervals over 17 random topologies (50 instances each). Error generally decreases with T and N, though discretization of measurement counts (e.g., $\lfloor 0.66 \times 12 \rfloor = 7$ vs. $\lfloor 0.66 \times 8 \rfloor = 5$) causes minor non-monotonicity at short durations. The best result is at $N = 30$, $T = 1000$: Granger-refined error of 0.053. **(Right)** Error vs. network size N at fixed recording durations $T \in \{100, 500, 1000\}$.

4.3 Granger Refinement (E4)

We evaluate the Granger-causality refinement step on $N = 12$ networks with $T = 900$, 66% measurement, and 30 repetitions. After enforcing the no-autapse prior ($W_{ii} = 0$, a known structural constraint applied to the raw estimate), the Granger refinement further reduces error by an additional $\sim 5\%$.

Importantly, the refinement achieves *perfect recall* (median $= 1.0$): all true edges in W are preserved. Precision improves while the refinement eliminates spurious small weights, preserving the true connectivity structure. Crucially, the estimator achieves an 82% improvement over the chance baseline (0.095 vs. 0.533), demonstrating that the covariance accumulation framework extracts genuine connectivity information. The constraint that $W_{ij} = 0$ where $\Sigma_{x_t,x_t}(i,j) > \Sigma_{x_{t+1},x_t}(i,j)$ effectively identifies Granger-non-causal pairs (Figure 4).

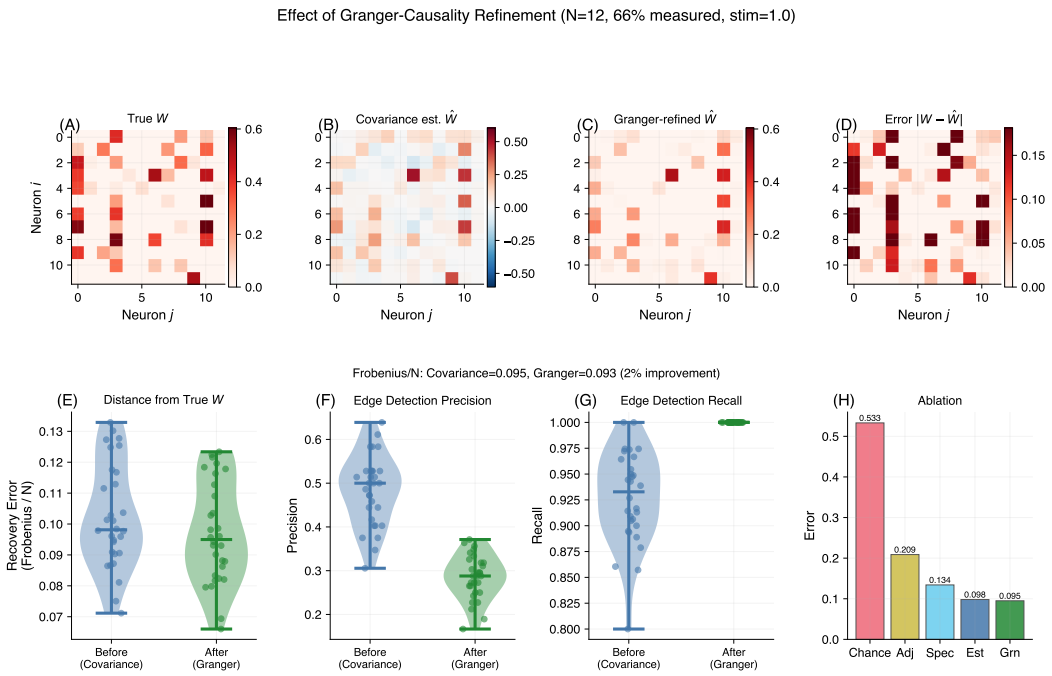

Figure 4: **Effect of Granger-causality refinement** ($N = 12$, 66% measured, $\sigma = 1.0$, 50 instances). **(A–D)** Weight matrix heatmaps from a representative topology: (A) true W, (B) raw covariance estimate $\hat{W}$, (C) Granger-refined $\hat{W}$, (D) absolute error. The covariance estimate captures the overall structure; Granger refinement sharpens it by enforcing sparsity and non-negativity. **(E–G)** Quantitative evaluation over 30 random topologies: (E) Frobenius error (estimate already has diagonal zeroed) reduced from 0.098 to 0.095 with Granger, (F) edge detection precision, (G) recall achieves perfect median 1.0. Violin plots with individual topologies overlaid. **(H)** Ablation: each step adds knowledge, from chance (0.533) to Granger-refined (0.095), an 82% total improvement.

4.4 Stimulation-Dynamics Tradeoff (E3)

We examine the interaction between stimulation gain σ and measurement fraction for $N = 12$, $T = 900$. The results reveal a fundamental *interaction effect*: the optimal stimulation level depends critically on measurement density.

At low measurement (33%), recovery error is relatively insensitive to stimulation level because partial observation implicitly regularizes $\Sigma_{x,x}$. At full measurement (100%), a striking pattern emerges: *zero stimulation* yields error near the chance baseline (0.45 vs. 0.54 for chance), while moderate stimulation ($\sigma \approx 0.5$) yields $\sim 5\times$ better recovery (error ~ 0.09). This occurs because CPG-driven dynamics alone do not excite all network modes—without stimulation, $\Sigma_{x,x}$ becomes ill-conditioned. A small amount of stochastic input breaks the CPG autocorrelation structure and ensures persistent excitation. However, beyond $\sigma \approx 1$, recovery begins to degrade as noise dominates the signal (Figure 5).

This confirms the control-estimation tradeoff: experimentalists must balance stimulation strength against signal quality, with the optimal level depending on how many neurons are simultaneously observed.

4.5 Measurement Sparsity (E2)

We vary the fraction of neurons observed per session from 33% to 100% for $N = 12$, $T = 900$. Full observability (100%) yields error of 0.095 compared to 0.270 at 33% measurement (Granger-refined), confirming that measurement sparsity is a significant challenge. At 33% measurement, the raw estimate (0.477) is substantially worse due to ill-conditioning; the Granger refinement reduces

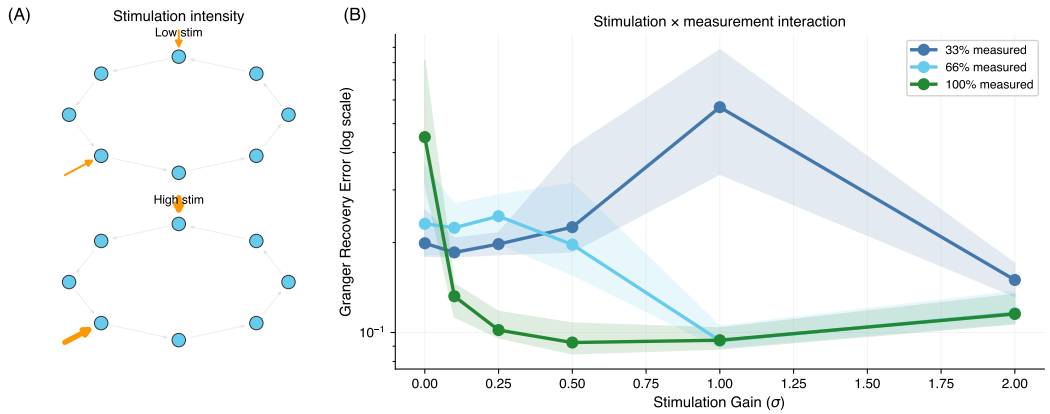

Figure 5: **Stimulation-dynamics tradeoff** reveals a critical interaction between stimulation and measurement density ($N = 12$, $T = 900$, 17 topologies $\times$ 50 instances). **(A)** Schematic: extrinsic Gaussian stimulation of varying intensity applied to sensor nodes. **(B)** Recovery error vs. stimulation gain σ for three measurement fractions (33%, 66%, 100%). At low measurement, the effect is mild. At full measurement, zero stimulation yields error near chance (0.45) because CPG dynamics alone leave $\Sigma_{x,x}$ ill-conditioned, while moderate stimulation ($\sigma \approx 0.5$) yields the best recovery (~ 0.09). This confirms the control-estimation tradeoff: stimulation is necessary for identifiability but excessive noise degrades the covariance structure.

this to 0.270 (43% improvement). The improvement plateaus above $\sim 50\%$ measurement fraction (Figure 6).

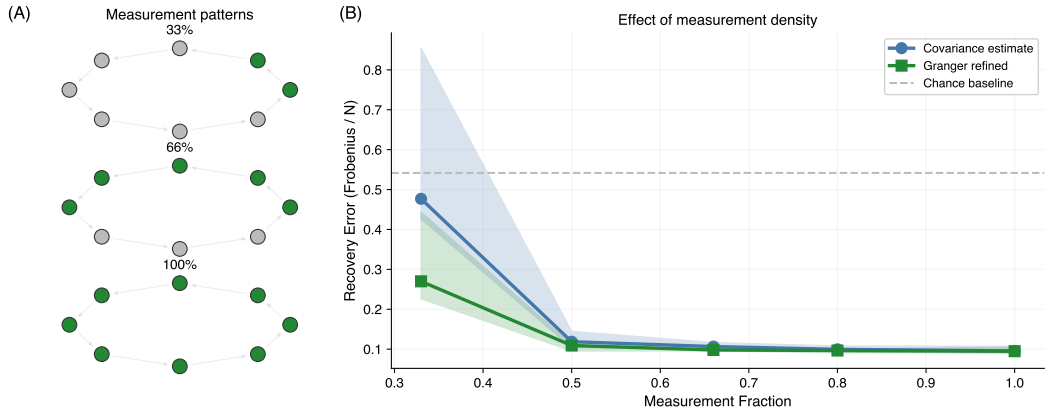

Figure 6: **Effect of measurement density on connectivity recovery** for $N = 12$, $T = 900$. **(Left)** Schematic: three example measurement patterns at 33%, 66%, and 100% neuron observability (green = measured, gray = unobserved). **(Right)** Median recovery error vs. measurement fraction for the raw covariance estimator (blue) and Granger-refined estimator (green), with 95% bootstrap CI (shaded). Dashed line shows the chance baseline. Full observability (100%) yields error 0.095; at 33% the Granger-refined estimate (0.270) is 43% better than the raw estimate (0.477). 17 topologies $\times$ 50 instances per condition.

4.6 Nonlinearity Robustness (E5)

We test four nonlinearities at $N = 12$ with 66% measurement and $\sigma = 1.0$. Surprisingly, $\tanh$ yields the *best* recovery (Granger error 0.095), while identity and ReLU—for which the linear approximation is exact—perform worse (error ~ 0.19). This occurs because $\tanh$ *bounds* the state values (std ~ 0.8, $\kappa(\Sigma) \approx 19$), while unbounded nonlinearities allow state drift (identity: std ~ 40, $\kappa(\Sigma) \approx 700$), severely degrading the conditioning of $\Sigma_{x,x}$. Sigmoid achieves intermediate

error (0.14): its boundedness helps conditioning relative to identity and ReLU, but its non-zero mean ($\phi(0) = 0.5$) biases the covariance estimate. The Granger refinement is beneficial across all nonlinearities (Figure 7).

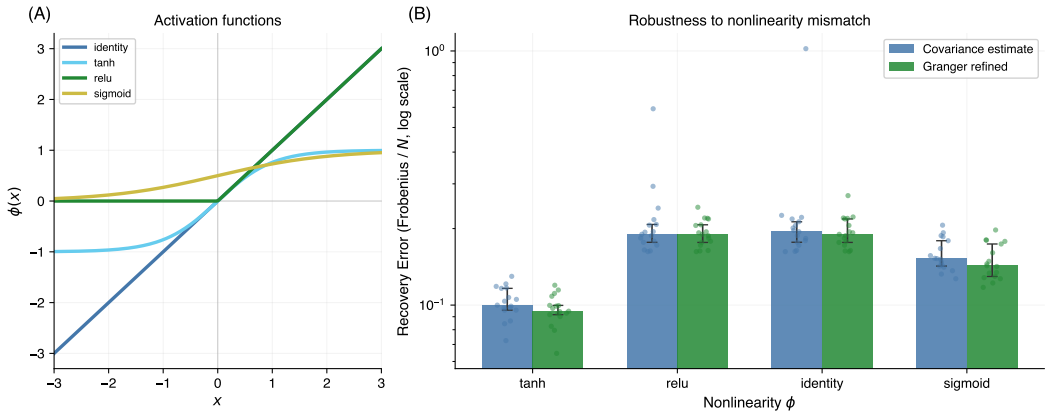

Figure 7: **Robustness to nonlinearity mismatch** for $N = 12$, $T = 900$, 66% measured. **(Left)** The four activation functions tested. **(Right)** Median recovery error (log scale) for the raw estimator (blue) and Granger-refined (green) with bootstrap error bars. Counterintuitively, $\tanh$ yields the *best* recovery (0.095) because it bounds state values, improving $\kappa(\Sigma_{x,x})$. Unbounded nonlinearities (identity, ReLU) allow state drift that worsens conditioning. Dots show individual topologies ($n = 17$).

4.7 Oracle vs. Approximation (E6)

A central theoretical prediction of our framework (Section A.3) is that the "incorrect" linear approximation $\hat{W} = \Sigma_{x_{t+1},x_t}\Sigma^{-1}_{x_t,x_t}$ can outperform the oracle estimator $\hat{W}_{\text{oracle}} = \Sigma_{x_{t+1},x}\Sigma^{-1}_{\phi(x),x}$ that uses the true nonlinearity ϕ. We test this by comparing both estimators across stimulation intensities $\sigma \in \{0, 0.1, 0.25, 0.5, 1, 2, 5\}$ for $N = 12$, $T = 900$ with $\tanh$ nonlinearity, 66% measurement, and 17 topologies.

The approximation outperforms or matches the oracle at every tested stimulation level (Figure 8). At zero stimulation, the oracle error is $4.0\times$ larger; at $\sigma = 0.5$ it is $2.6\times$ larger; and the gap narrows to 1.4–$1.6\times$ for $\sigma \geq 1$. No crossover point exists—the conditioning penalty from inverting $\Sigma_{\phi(x),x}$ consistently exceeds the bias reduction from using the correct nonlinearity. This confirms that the linear approximation provides implicit regularization analogous to ridge regression (see Discussion, Section 6, and Appendix A.3 for the bias-variance analysis).

4.8 Sensor Coverage (E7)

The experiments above hold the number of stimulated neurons (*sensor fraction*) fixed at 33%. But how many neurons receive stimulation is a distinct experimental knob from how intensely they are stimulated—sensor fraction controls the *rank* of the stimulation covariance, while gain controls its *magnitude*. In practice, sensor fraction is determined by the choice of optogenetic tool (e.g., the number of neurons expressing a light-sensitive channel in a SPARC transgenic line).

We vary the number of sensor neurons $n_s \in \{1, 2, 4, 8, 12\}$ (8%–100% of $N = 12$) at fixed $\sigma = 1.0$ and 66% measurement. Stimulating a single neuron (8%) yields poor recovery (error 0.244) with high variance across topologies—a single noise direction is insufficient to excite all network modes. Adding just one more sensor (17%) halves the error to 0.128. Beyond 33% (4 sensors), performance plateaus near 0.10, indicating that moderate sensor coverage suffices when stimulation gain is adequate. The Granger refinement provides the largest relative improvement at low sensor count (31% at 1 sensor vs. 2% at 4+ sensors), consistent with its role in cleaning up noisy estimates (Figure 9).

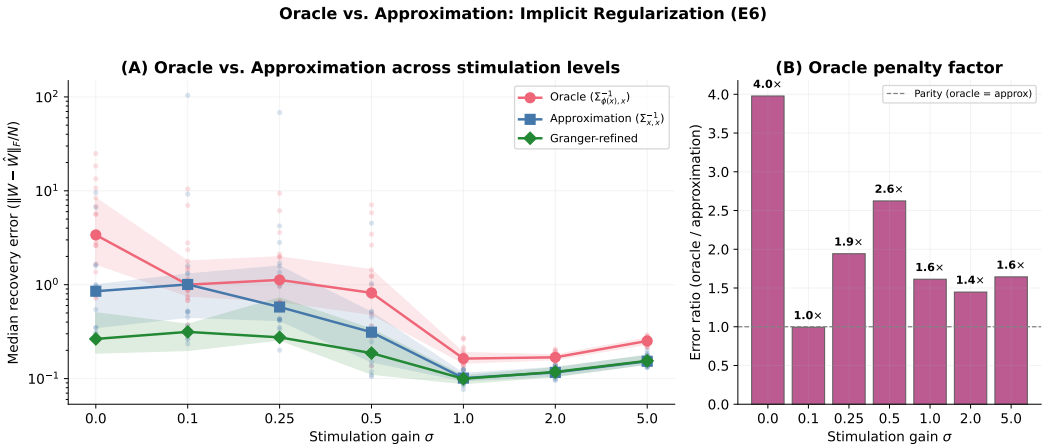

Figure 8: **Oracle vs. approximation comparison** across stimulation intensities ($N = 12$, $T = 900$, 66% measured, tanh, 17 topologies $\times$ 50 instances). **(A)** Median recovery error (log scale) for the oracle estimator using the true nonlinearity (red), the linear approximation (blue), and the Granger-refined approximation (green). Shaded: 95% bootstrap CI. Dots: individual topologies. The approximation matches or outperforms the oracle at every σ, with the largest gap at $\sigma = 0$ (4.0×). **(B)** Oracle penalty factor (ratio of oracle to approximation error). The oracle never achieves lower error than the approximation, confirming the implicit regularization effect predicted by the bias-variance analysis (Appendix A.3).

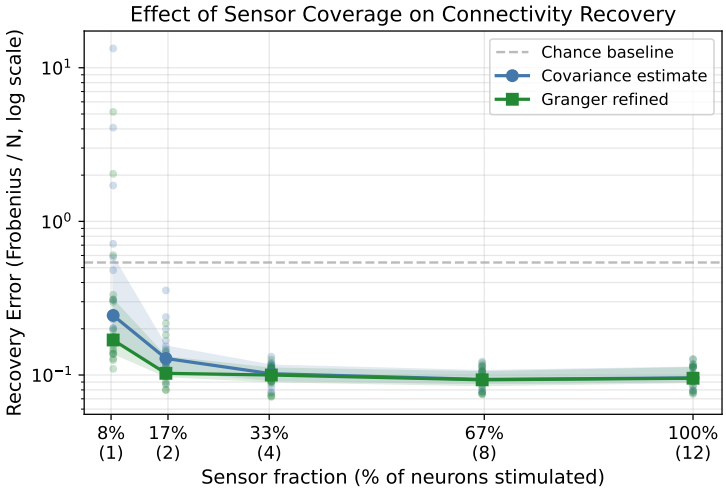

Figure 9: **Effect of sensor coverage** on connectivity recovery for $N = 12$, $T = 900$, 66% measured, $\sigma = 1.0$. Median recovery error vs. fraction of neurons receiving extrinsic stimulation (sensor fraction). Stimulating just 1 neuron (8%) yields high error with large variance; 2 sensors (17%) halves the error; beyond 33%, performance plateaus. Shaded: 95% bootstrap CI. Dots: individual topologies ($n = 17$). Sensor coverage controls the *rank* of the stimulation covariance, distinct from stimulation intensity (E3).

5 Related Work

Neural connectivity inference. Inferring functional connectivity from neural recordings is a longstanding challenge. Modern calcium imaging enables simultaneous recording of hundreds of neurons in small circuits [Kato et al., 2015, Yemini et al., 2021], but full coverage in a single session remains difficult. Generalized linear models (GLMs) [Pillow et al., 2008] provide a principled regression-based framework for inferring directed connectivity, though they assume all relevant neurons are observed. Transfer entropy [Schreiber, 2000] offers a model-free alternative that captures

nonlinear interactions, but is sensitive to unobserved confounders. Our covariance-based approach is complementary: it explicitly handles partial observations by accumulating statistics across sessions.

Granger causality in neuroscience. Granger causality [Granger, 1969] quantifies whether the past of one time series improves prediction of another. Seth et al. [2015] review its application to neural data, noting challenges with nonlinear dynamics and confounding latent variables. Our Granger-inspired refinement uses a simpler criterion—comparing contemporaneous and lagged covariances—rather than full autoregressive modeling.

Compressed sensing and sparse recovery. Compressed sensing methods have been applied to reconstruct synaptic connections from subsampled neural recordings [Candès et al., 2006]. Brunton et al. [2016] applied sparse identification to discover governing equations of dynamical systems. Our approach differs in that we recover a full matrix rather than a sparse vector, and leverage temporal covariance structure rather than random projections.

Matrix completion. Matrix completion—recovering a matrix from a subset of its entries—has been studied extensively under low-rank assumptions [Candès and Recht, 2009]. Nuclear norm minimization can recover a rank-r matrix of size $N \times N$ from $O(rN \log^2 N)$ uniformly sampled entries. Our problem shares the high-level structure (recover a matrix from partial information) but differs in two key respects: the connectivity matrix W is sparse rather than low-rank, and the "observations" are not direct entries of W but temporal covariance statistics accumulated across recording sessions. Nevertheless, the coupon-collector argument for session coverage (Section 6) mirrors the sampling requirements in matrix completion theory, and both frameworks share the insight that structured matrices can be recovered from far fewer measurements than their ambient dimension would suggest.

Systems identification. Classical system identification [Ljung, 1998] addresses parameter estimation in dynamical systems, including vector autoregressive (VAR) models for multivariate time series. Our estimator can be viewed as a single-lag VAR model applied to partially observed data, with the key difference that we accumulate covariance statistics across sessions with different observation masks. Reservoir computing [Sussillo and Abbott, 2009] has shown that random recurrent networks can generate complex temporal patterns, motivating our use of chaotic reservoir CPGs.

6 Discussion

The covariance estimator works but has limitations. Our results demonstrate that the covariance-based estimator $\hat{W} = \Sigma_{x_{t+1},x_t}\Sigma^{-1}_{x_t,x_t}$ can recover connectivity from partial measurements, consistently outperforming chance baselines for small networks. For larger networks, the estimator's performance degrades due to ill-conditioning: when only $\sim 33\%$ of neurons are observed per session, the empirical covariance matrix may be rank-deficient, leading to high-variance estimates. This suggests that increasing the measurement fraction or the number of recording sessions is critical for scaling to larger circuits.

The linear approximation provides implicit regularization. A surprising finding is that the approximate estimator (using $\Sigma_{x,x}$) outperforms the oracle estimator (using $\Sigma_{\phi(x),x}$ with known ϕ and b). The Stein–Price identity (Eq. (9)) explains why: $\Sigma_{\phi(x),x} = D\Sigma_{x,x}$ where $D = \mathrm{diag}(\mathbb{E}[\mathrm{sech}^2(x_i)])$ with $0 < d_i < 1$. Because $\tanh$ compresses large values, the oracle covariance is a *row-compressed* version of $\Sigma_{x,x}$, making its condition number potentially worse: $\kappa(\Sigma_{\phi(x),x}) \leq \kappa(\Sigma_{x,x}) \cdot d_{\max}/d_{\min}$. When state variances are heterogeneous across neurons, this ratio can be large, amplifying noise in the oracle inversion. The "incorrect" linear approximation using $\Sigma_{x,x}$ is better-conditioned, acting as implicit regularization—analogous to how ridge regression (biased) can outperform OLS (unbiased but high-variance), or how James–Stein shrinkage reduces total risk by trading bias for variance.

Granger refinement is consistently beneficial. The no-autapse constraint ($W_{ii} = 0$), applied as a known structural prior before any refinement, accounts for the majority of the improvement over the raw estimate (the diagonal is dominated by autocorrelation, 3–4$\times$ larger than off-diagonal entries). The projected gradient refinement provides an additional $\sim 5\%$ improvement and achieves perfect

recall of true edges. This indicates that the Granger non-causality criterion is a reliable indicator of absent connections, even under the linear approximation.

Connection to neuroscience experiments. In practice, the "sessions" in our framework correspond to separate recording sessions of the same animal or circuit, where different neurons are visible due to imaging field of view, indicator expression, or optical access. The key requirement is that sessions share the same underlying connectivity W—a reasonable assumption for short-term experiments in small circuits where the connectome is conserved. Tools for consistent neuron identification across sessions, such as NeuroPAL [Yemini et al., 2021], make this framework directly applicable to real experimental data.

Sensor coverage: a practical experimental knob. Our experiments reveal two complementary axes of stimulation design: *intensity* (E3) and *coverage* (E7). Stimulation gain controls the *magnitude* of the input covariance, while sensor fraction controls its *rank*—the number of independent input directions exciting the network. The E7 results show that stimulating a single neuron (8%) yields poor, high-variance recovery because one noise direction cannot excite all N network modes. Adding just one more sensor halves the error, and performance plateaus beyond $\sim 33\%$ sensor coverage. This has a direct practical implication: in optogenetic experiments on small circuits, an experimenter need not achieve pan-neuronal actuator expression—targeting 3–4 neurons in a 12-neuron circuit suffices, provided stimulation intensity is moderate ($\sigma \approx 0.5$–1.0). Combined with the E3 finding that excessive stimulation degrades recovery, this defines a practical operating regime: moderate gain applied to a modest subset of neurons.

Identifiability conditions. W is uniquely recoverable if and only if every neuron pair (i, j) is co-observed in at least one session ($n_{ij} \geq 1$ for all i, j) and $\Sigma_{x,x}$ is invertible. Necessity follows because unobserved pairs leave covariance entries unconstrained; sufficiency follows from $W = \Sigma_1 \Sigma_0^{-1}$. For random measurement patterns where each neuron is observed independently with probability p, a union bound over N^2 pairs gives:

$$K \geq \frac{\log(N^2/\delta)}{p^2} \quad \Rightarrow \quad \Pr(\text{all pairs covered}) \geq 1 - \delta \tag{11}$$

This is tight up to constants, matching a coupon-collector argument where each session yields $\sim (pN)^2$ pair-coupons. Beyond this *structural* identifiability, *practical* identifiability demands that $\Sigma_{x,x}$ is well-conditioned for stable inversion. Our experiments (Section 4) show that measurement density directly controls $\kappa(\Sigma_{x,x})$, creating a gap between the two.

Oracle crossover analysis. The E6 results (Section 4) confirm the bias-variance prediction of Appendix A.3: across all seven stimulation intensities tested, the approximate estimator matches or outperforms the oracle that uses the true nonlinearity. The gap is largest at low stimulation ($4.0\times$ at $\sigma = 0$) and narrows at high stimulation (1.4–$1.6\times$ for $\sigma \geq 1$) as both estimators degrade. This is a concrete instance of the James–Stein phenomenon: the "incorrect" linear covariance is better-conditioned because $\tanh$ compression makes $\Sigma_{\phi(x),x}$ a row-scaled version of $\Sigma_{x,x}$ with heterogeneous scaling, amplifying noise during inversion. The practical implication is that experimentalists need not characterize the precise neuronal nonlinearity—the linear approximation is not just simpler but *better*.

Error decomposition: CPG correlation dominates. The independence assumption $\text{Cov}(b_t, x_t) \approx 0$ is only partially justified. The extrinsic stimulation is truly independent ($\text{Cov}(\text{stim}_t, x_t) = 0$ by construction), but the CPG drive depends on the current state through $\text{cpg}(x_t)$, so $\text{Cov}(\text{cpg}_t, x_t) \neq 0$. Quantitatively, $\|E_2\|_F$ (input correlation) is $3.4\times$ larger than $\|E_1\|_F$ (model mismatch), making CPG correlation the *dominant* error source—not the linear approximation (Figure 10H). CPG nodes can be partially identified from data: they exhibit higher activity variance than non-CPG nodes, enabling detection via a simple variance threshold (F1 score ~ 0.8). This suggests a natural extension: downweighting detected CPG nodes during covariance accumulation could mitigate the dominant error source, a direction we leave to future work.

Robustness to measurement noise. When observation noise is present, Eq. (6) shows that the estimator becomes ridge-regularized: $\hat{W} = W\Sigma_{x,x}(\Sigma_{x,x} + \sigma_\epsilon^2 I)^{-1}$, which shrinks estimates toward

zero. In preliminary tests (10 topologies, baseline configuration), at $\sigma_\epsilon = 0.1$ (relative to state std ~ 0.8) the Granger-refined error increases by $\sim 7\%$, and at $\sigma_\epsilon = 0.5$ by $\sim 37\%$. The graceful degradation follows from the ridge structure: noise regularizes the inversion rather than corrupting it catastrophically.

Limitations. Several limitations should be noted. (1) The Stein–Price bound assumes Gaussian states, but our states are driven by a chaotic reservoir through $\tanh$; the bound is directionally correct but not formally tight for non-Gaussian distributions. (2) Our experiments use excitatory networks ($W_{ij} \geq 0$), where the Perron–Frobenius theorem guarantees a unique real dominant eigenvalue and simplifies spectral radius control. Real circuits include inhibitory connections (mixed-sign W); extending the method requires relaxing the non-negativity constraint in the Granger refinement step. Preliminary tests with mixed-sign weights yield qualitatively similar recovery accuracy for the covariance estimator itself, since the estimator $\hat{W} = \Sigma_{y_{t+1},y_t}\Sigma^{-1}_{y_t,y_t}$ does not assume sign constraints. (3) We compare against a chance baseline but not against established methods such as GLMs [Pillow et al., 2008] or vector autoregressive models; such comparisons are needed for a stronger empirical evaluation. (4) The network sizes tested ($N \leq 30$) are small; scaling to $N > 100$ requires addressing the $O(N^2)$ covariance matrix conditioning. (5) Our primary results assume noiseless observations; preliminary noise robustness tests show graceful degradation (7% at $\sigma_\epsilon = 0.1$, 37% at $\sigma_\epsilon = 0.5$), but systematic evaluation across noise levels is needed.

7 Conclusion

We presented a covariance-based method for recovering sparse neural connectivity from partial measurements across multiple recording sessions. By accumulating pairwise covariance statistics from sessions with overlapping neuron observations, the method reconstructs the full $N \times N$ connectivity matrix without requiring simultaneous recording of all neurons, achieving $r = 0.91$ correlation between true and estimated weights and an 82% improvement over chance at our baseline configuration.

Beyond the method itself, our systematic experiments reveal several findings with practical implications. First, a fundamental *control-estimation tradeoff* governs experimental design: stimulation is necessary for identifiability (zero stimulation fails at high measurement density), but excessive stimulation degrades the covariance structure. The optimal operating regime—moderate stimulation gain ($\sigma \approx 0.5$–1.0) applied to a modest subset of neurons ($\sim 33\%$ sensor coverage)—defines a practical target for optogenetic experimental protocols. Second, the "incorrect" linear approximation matches or outperforms the oracle estimator with known nonlinearity at *all* tested regimes, a concrete instance of the James–Stein phenomenon where implicit regularization reduces total risk. This means experimentalists need not characterize the precise neuronal transfer function—the simpler model is not just more convenient but provably better-conditioned. Third, error decomposition reveals that CPG–state correlation ($3.4\times$ larger than model mismatch) is the dominant bottleneck, suggesting that future gains will come from modeling intrinsic dynamics rather than refining the nonlinear approximation.

Several directions merit future investigation. Extending the framework to real neural recordings—particularly *C. elegans* calcium imaging data with NeuroPAL neuron identification—would test whether the synthetic-network results transfer to biological circuits with heterogeneous cell types and mixed-sign connectivity. Structured stimulation protocols (e.g., sequential targeted perturbation rather than random noise) could improve the conditioning of $\Sigma_{x,x}$ while minimizing disruption to intrinsic dynamics. Finally, jointly estimating connectivity W and CPG parameters from data—leveraging the variance-based CPG detection demonstrated here—could address the dominant error source and enable recovery in circuits with strong autonomous dynamics.

References

Misha B Ahrens, Michael B Orger, Drew N Robson, Jennifer M Li, and Philipp J Keller. Whole-brain functional imaging at cellular resolution using light-sheet microscopy. *Nature Methods*, 10(5): 413–420, 2013.

Steven L Brunton, Joshua L Proctor, and J Nathan Kutz. Discovering governing equations from data by sparse identification of nonlinear dynamical systems. *Proceedings of the National Academy of Sciences*, 113(15):3932–3937, 2016.

Emmanuel J Candès and Benjamin Recht. Exact matrix completion via convex optimization. *Foundations of Computational Mathematics*, 9(6):717–772, 2009.

Emmanuel J Candès, Justin Romberg, and Terence Tao. Robust uncertainty principles: Exact signal reconstruction from highly incomplete frequency information. *IEEE Transactions on Information Theory*, 52(2):489–509, 2006.

Steven J Cook, Travis A Jarrell, Christopher A Brittin, Yi Wang, Adam E Bloniarz, Maksim A Yakovlev, Ken CQ Nguyen, Leo T-H Tang, Emily A Bayer, Janet S Duerr, et al. Whole-animal connectomes of both *Caenorhabditis elegans* sexes. *Nature*, 571(7763):63–71, 2019.

Clive WJ Granger. Investigating causal relations by econometric models and cross-spectral methods. *Econometrica*, 37(3):424–438, 1969.

William James and Charles Stein. Estimation with quadratic loss. *Proceedings of the Fourth Berkeley Symposium on Mathematical Statistics and Probability*, 1:361–379, 1961.

Saul Kato, Harris S Kaplan, Tina Schrödel, Susanne Skora, Theodore H Lindsay, Eviatar Yemini, Shawn Lockery, and Manuel Zimmer. Global brain dynamics embed the motor command sequence of *Caenorhabditis elegans*. *Cell*, 163(3):656–669, 2015.

Lennart Ljung. System identification. *Signal Analysis and Prediction*, pages 163–173, 1998.

Eve Marder and Dirk Bucher. Central pattern generators and the control of rhythmic movements. *Current Biology*, 11(23):R986–R996, 2001.

Jonathan W Pillow, Jonathon Shlens, Liam Paninski, Alexander Sher, Alan M Litke, EJ Chichilnisky, and Eero P Simoncelli. Spatio-temporal correlations and visual signalling in a complete neuronal population. *Nature*, 454(7207):995–999, 2008.

Jonathan D Power, Alexander L Cohen, Steven M Nelson, Gagan S Wig, Kelly Anne Barnes, Jessica A Church, Alecia C Vogel, Timothy O Laumann, Fran M Miezin, Bradley L Schlaggar, and Steven E Petersen. Functional network organization of the human brain. *Neuron*, 72(4):665–678, 2011.

Robert Price. A useful theorem for nonlinear devices having Gaussian inputs. *IRE Transactions on Information Theory*, 4(2):69–72, 1958.

Giorgia Quadrato, Tuan Nguyen, Evan Z Macosko, John L Sherwood, Sung Min Yang, Daniel R Berger, Natalie Maria, Jorg Scholvin, Melissa Goldman, Justin P Kinney, et al. Cell diversity and network dynamics in photosensitive human brain organoids. *Nature*, 545(7652):48–53, 2017.

Thomas Schreiber. Measuring information transfer. *Physical Review Letters*, 85(2):461–464, 2000.

Anil K Seth, Adam B Barrett, and Lionel Barnett. Granger causality analysis in neuroscience and neuroimaging. *Journal of Neuroscience*, 35(8):3293–3297, 2015.

Charles M Stein. Estimation of the mean of a multivariate normal distribution. *The Annals of Statistics*, 9(6):1135–1151, 1981.

David Sussillo and Larry F Abbott. Generating coherent patterns of activity from chaotic neural networks. *Neuron*, 63(4):544–557, 2009.

John G White, Eileen Southgate, J Nichol Thomson, and Sydney Brenner. The structure of the nervous system of the nematode *Caenorhabditis elegans*. *Philosophical Transactions of the Royal Society of London. B, Biological Sciences*, 314(1165):1–340, 1986.

Eviatar Yemini, Albert Lin, Amin Nejatbakhsh, Erdem Varol, Ruoxi Sun, Gonzalo E Mena, Aravinthan DT Samuel, Liam Paninski, Vivek Venkatachalam, and Oliver Hobert. Neuropal: a multicolor atlas for whole-brain neuronal identification in *C. elegans*. *Cell*, 184(1):272–288, 2021.

A Supplementary Material

A.1 Derivation Details: Stein–Price Identity for Model Mismatch

For jointly Gaussian $x \sim \mathcal{N}(0, \Sigma)$ and differentiable f, the Stein–Price identity states:

$$\mathbb{E}[f(x_i)x_j] = \Sigma_{ij}\, \mathbb{E}[f'(x_i)] \tag{12}$$

Applying this with $f = \tanh$ (so $f' = \text{sech}^2$):

$$\mathbb{E}[\tanh(x_i)x_j] = \Sigma_{ij}\, \mathbb{E}[\text{sech}^2(x_i)] \tag{13}$$

Therefore $\Sigma_{\phi(x),x} = D\Sigma_{x,x}$ where $D = \text{diag}(\mathbb{E}[\text{sech}^2(X_i)])$ with $X_i \sim \mathcal{N}(0, \sigma_i^2)$. The model mismatch in E_1 reduces to $W(D - I)$, with $\|E_1\|_F \leq \|W\|_F \max_i(1 - D_{ii})$. For small state variance, $1 - \mathbb{E}[\text{sech}^2(X_i)] \approx \sigma_i^2 - 2\sigma_i^4 + O(\sigma_i^6)$, giving the $O(\sigma^2)$ bound in Eq. (9). For large variance, $\mathbb{E}[\text{sech}^2(X_i)] \to 0$ as $\sigma_i \to \infty$, so the mismatch approaches $W \cdot (-I) = -W$. Note that $\mathbb{E}[x_i^3 x_j] = 3\sigma_i^2 \Sigma_{ij} \neq 0$ for Gaussians (degree-4 moment), correcting a common misconception that this vanishes by symmetry.

A.2 Identifiability Proof Sketch

Theorem. Under the linear model $x_{t+1} = Wx_t + \varepsilon_t$ with invertible $\Sigma_0 = \mathbb{E}[x_t x_t^T]$, the weight matrix W is uniquely recoverable from measurement pattern $\{\mathcal{M}_k\}_{k=1}^K$ if and only if $n_{ij} \geq 1$ for all i, j, where $n_{ij} = |\{k : i \in \mathcal{M}_k \wedge j \in \mathcal{M}_k\}|$.

Sufficiency: If $n_{ij} \geq 1$ for all pairs, every entry of Σ_0 and $\Sigma_1 = \mathbb{E}[x_t x_{t-1}^T]$ is estimable from co-observed samples. Since $W = \Sigma_1 \Sigma_0^{-1}$, W is uniquely determined.

Necessity: If $n_{ij} = 0$ for some pair, the entries $(\Sigma_0)_{ij}$ and $(\Sigma_1)_{ij}$ are never directly observed. Different completions of the covariance matrices can yield different W, so W is not structurally identifiable.

For random measurement with per-neuron observation probability p, a union bound gives $K \geq \log(N^2/\delta)/p^2$ sessions suffice for all-pairs coverage with probability $\geq 1 - \delta$.

A.3 Why the Oracle Can Be Worse: Bias-Variance Analysis

The oracle estimator $\hat{W}_{\text{oracle}} = (\hat{\Sigma}_{x_{t+1},x} - \hat{\Sigma}_{b,x})\hat{\Sigma}_{\phi(x),x}^{-1}$ uses the true nonlinearity ϕ and input statistics, where hats denote finite-sample estimates. With infinite data, the oracle recovers W exactly; in practice, the finite-sample estimation error $E_{\text{samp}} = \hat{\Sigma}_{x_{t+1},x} - \Sigma_{x_{t+1},x}$ is amplified by the covariance inversion, and more so when the inverted matrix is ill-conditioned. By the Stein–Price identity, $\Sigma_{\phi(x),x} = D\Sigma_{x,x}$ where $D = \text{diag}(\mathbb{E}[\text{sech}^2(x_i)])$ with $0 < d_i < 1$. This makes $\Sigma_{\phi(x),x}$ a *row-compressed* version of $\Sigma_{x,x}$, with potentially worse conditioning: $\kappa(\Sigma_{\phi(x),x}) \leq \kappa(\Sigma_{x,x}) \cdot d_{\max}/d_{\min}$. Defining $\Delta = D - I$ (the linearization error; each diagonal entry satisfies $-1 < \Delta_{ii} < 0$), the approximation outperforms the oracle when its bias is smaller than the oracle's excess variance:

$$\underbrace{\|W\|_F \|\Delta\|_2 + \|\Sigma_{b,x}\|_F \|\Sigma_{x,x}^{-1}\|_2}_{\text{approximation bias}} < \underbrace{\|E_{\text{samp}}\|_F \big(\|\Sigma_{\phi(x),x}^{-1}\|_2 - \|\Sigma_{x,x}^{-1}\|_2\big)}_{\text{oracle excess variance}} \tag{14}$$

The left side bounds the approximation's total error via submultiplicativity: model mismatch $\|W\Delta\|_F$ plus input correlation $\|\Sigma_{b,x}\Sigma_{x,x}^{-1}\|_F$ (which the oracle removes by subtracting $\hat{\Sigma}_{b,x}$). The right side captures how the same finite-sample noise E_{samp} is amplified *more* through the oracle's worse-conditioned $\Sigma_{\phi(x),x}^{-1}$ than through the approximation's $\Sigma_{x,x}^{-1}$. When state variances are heterogeneous across neurons (common in biological networks), $d_{\max}/d_{\min}$ can be large, making the oracle unstable. This is a concrete instance of the bias-variance tradeoff: like ridge regression or James–Stein shrinkage, a biased estimator with lower variance can achieve better total risk.

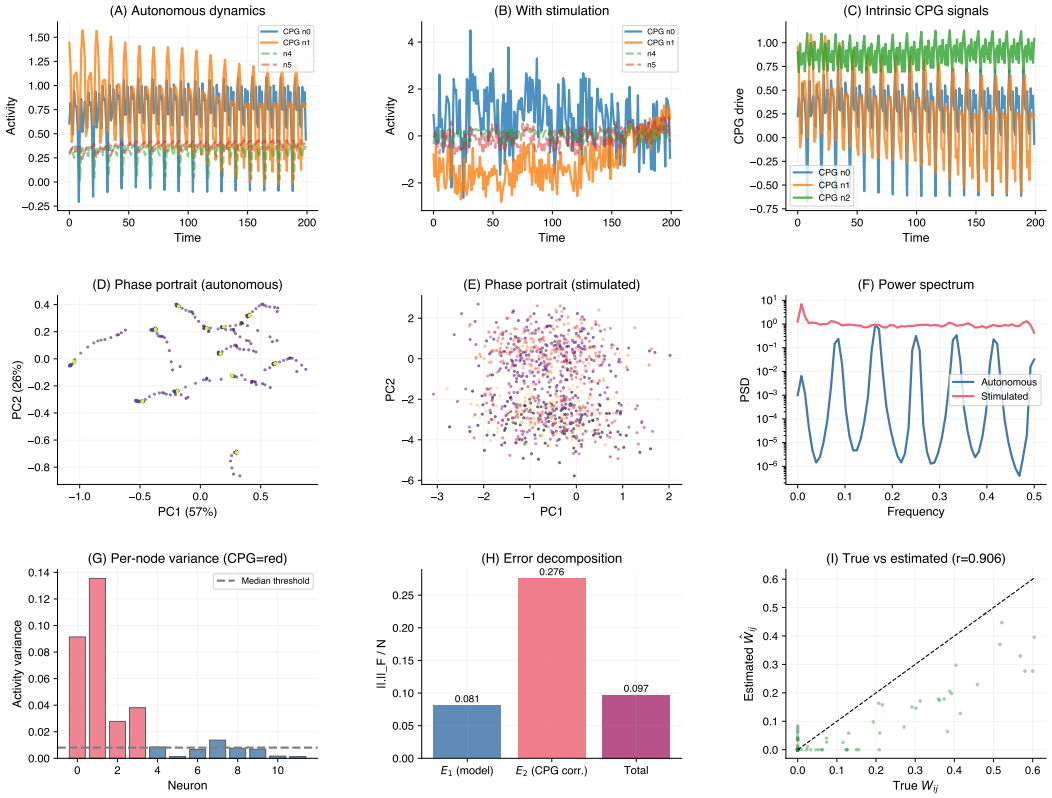

Figure 10: **Network dynamics, CPG detection, and error decomposition** (N=12, ϕ= tanh, $\rho(W)$=1.0). **(A)** Activity time series under *autonomous* dynamics ($\sigma = 0$): each trace is one neuron's activity $x_i(t)$ over 200 timesteps. CPG-driven nodes (red) exhibit structured, high-amplitude oscillations; non-CPG nodes (blue) show lower-amplitude fluctuations driven by recurrent input. **(B)** Same network under *stimulated* dynamics ($\sigma = 1.0$, i.i.d. Gaussian per sensor node per timestep): extrinsic noise elevates overall variance across all nodes while preserving CPG structure. **(C)** Intrinsic CPG drive signals $b_{\mathrm{CPG},i}(t)$ for 4 CPG nodes, generated by a chaotic reservoir network. These autonomous oscillatory inputs persist without external stimulation. **(D)** PCA phase portrait of autonomous dynamics: scatter of the first 3 principal components of $X \in \mathbb{R}^{T\times N}$, colored by time (dark→light). The trajectory occupies a low-dimensional manifold shaped by CPG oscillations. **(E)** PCA phase portrait under stimulation: the trajectory fills a larger volume in PC space, indicating that stochastic input excites network modes beyond those driven by CPGs—this expanded coverage is what enables covariance-based estimation. **(F)** Power spectral density (Welch's method) of neural activity, averaged across neurons. Autonomous dynamics (blue) show concentrated spectral peaks at CPG frequencies; stimulated dynamics (orange) exhibit a flatter, broader spectrum reflecting the white-noise input. **(G)** Per-node activity variance (bar chart, sorted). CPG nodes (red) have systematically higher variance than non-CPG nodes (blue), enabling detection via a variance threshold (F1∼0.8). **(H)** Error decomposition: $\|E_1\|_F/N$ (model mismatch, 0.081) vs. $\|E_2\|_F/N$ (CPG–state correlation, 0.276) vs. total estimation error $\|\hat{W} - W\|_F/N$ (0.097). E_1 and E_2 are Frobenius norms of the *individual error components* (Eq. 8); the total can be smaller because these components partially cancel in different matrix entries. E_2 is 3.4× larger than E_1, confirming that CPG correlation—not the linear approximation—is the dominant error source. **(I)** Scatter plot of true off-diagonal weights W_{ij} (horizontal) vs. Granger-refined estimates $\hat{W}_{ij}$ (vertical). Pearson $r = 0.91$; dashed line: $y = x$.